

Real-Time Multi-Domain Big Data Analytics for Smart Cities: Integration, Privacy, and Performance Evaluation

Mannat

*Chitkara University Institute of Engineering and Technology, Chitkara University, Punjab
India*

mannat557.be22@chitkara.edu.in

Abstract—The complexity of managing contemporary cities has increased due to their rapid growth. In an effort to address these issues, smart city initiatives use sensors and computers to monitor and enhance municipal services. An essential tool for smart cities is big data analytics. It facilitates the simultaneous processing of numerous types of data. Roads, power grids, and safety systems are examples of city systems that constantly produce large amounts of data that require careful analysis. This study focuses on a data analytics system that can evaluate large amounts of city data in real time and provide practical suggestions for improving city services.

Combining all the data, protecting the data, and improving the analytics system are the three main goals of the system. The system can analyse a large amount of data from Internet of Things devices and city systems using computers and smart analytics techniques. The study also evaluates the system's ability to handle large amounts of data, process data quickly, and assist users in making decisions. The system uses techniques like limiting who can view the data and concealing who the data is about in order to protect city data. The findings demonstrate how real-time data analytics can improve city services by facilitating quicker decision-making and more efficient use of resources.

Methodology used in this study has a positive contribution towards building a secure analytics system that will grow together with the city and hence foster sustainable growth within the city thus ensuring quality life for the city dwellers. There are several benefits that can be realized through big data analytics when it comes to cities. To start with, there will be enhanced effectiveness and efficiency in service delivery within the city. Since big data analytics helps city officials to improve their decision making and service delivery, this kind of a system is necessary in cities. It is important to note that big data analytics will play a role in helping to make the city more livable and sustainable. In conclusion, this study highlights how big data analytics is instrumental in improving city operations.

Keywords—Smart Cities, Big Data Analytics, Real-Time Data Processing, IoT, Urban Data Integration, Data Privacy, Distributed Computing, Urban Infrastructure Optimization

I. INTRODUCTION

Urbanization is changing cities fast. Now most of the world's population lives in cities. This number will go up a lot in the coming years. As cities get bigger it becomes harder to manage things, like energy distribution, public safety, healthcare, transportation networks and keeping an eye on the environment. The old ways of managing cities, which use data and people making decisions by hand are not working for today's cities because they are so big and always changing. Urbanization

and cities are a problem. Managing cities is a job because of urbanization.

The concept of smart cities is really taking off as a way to deal with the problems that cities face. Smart cities make use of things like the Internet of Things, cloud computing, and big data analytics to keep an eye on what's going on, look at the information, and make urban operations better. There is a lot of data coming in all the time from sensors, devices that are connected to the internet, and smart systems. This data includes things like what's happening with traffic, how energy is being used, what the environment is like, and statistics on public safety. Looking at this smart city data is very important if you want to make decisions, use resources in a smart way, and make the lives of the people who live in smart cities better. Smart cities are about using this kind of technology to make life easier for everyone.

However, there are many difficulties due to the quick expansion of data in smart city ecosystems. The high volume, speed, and variety of the data make it challenging to process using conventional techniques. Real-time responses are also necessary for many urban services, such as energy balance, emergency response, and traffic management. Even though big data analytics can handle such complicated datasets, current solutions frequently have trouble handling real-time multi-domain data integration. Furthermore, concerns about data security, privacy, and interoperability between various urban systems are not sufficiently addressed.

A. Objective of the Work

- We want to create a system that can look at data from many areas in a city and analyze it in real time for smart cities.
- We need to put all the different kinds of information from around the city into one place.
- We have to make sure that the people's private information is safe and that we are handling the data in a safe way.
- We want to make the system faster and more efficient so it can handle a lot of data and give us answers quickly.
- We hope this will help the people, in charge make decisions and make the city a better place to live.

By giving the city a system that can handle a lot of data and keep it safe we can help the city make decisions and make

the services they provide to people more efficient. This project is helping to create cities that are good for the environment, can withstand problems and use data to make good decisions, which makes the city a better place for people to live.

II. LITERATURE REVIEW

Smart cities have gained much interest because of urbanization and efficient utilization of resources. Big data analytics plays an important role in analyzing urban data in large quantities and improving operations within the city.

Several authors have investigated the usage of big data analysis within smart cities applications. The authors Batty *et al.* [1] stressed the importance of using a data-driven approach in urban planning and pointed out the importance of including digital technologies within urban infrastructures. Chen *et al.* [2] have reviewed various applications of big data technologies used in large-scale data analysis.

For the transport sector, research revealed the possibility of using GPS devices and sensors to collect real-time traffic data, which can be analyzed to enhance traffic management by reducing travel time [3]. For the energy sector, applications of big data and smart grids were used in the management of energy supply and demand in order to reduce energy wastage [4].

Furthermore, big data analytics has been used in waste management, where sensors help analyze real-time data for efficient waste collection routing [5]. Public safety has benefited from the use of analytics in processing data from surveillance systems and social media to support emergency response operations and high-risk area identification [6].

Both cloud computing and edge computing are important aspects when considering support for smart city analytics. The former provides scalable storage and processing services, while the latter decreases latency by processing data close to its source [7]. In addition, machine learning algorithms are frequently used in smart cities for predicting various events such as traffic flows, energy usage, and environmental monitoring [8].

While these approaches provide many benefits for smart city analytics, current research tends to focus on one domain at a time. There have not been many studies focusing on real-time multi-domain data integration, which also presents challenges related to data security and privacy.

Limitations of Existing Work

- Most research concentrates on applications in single domain instead.
- Low support of real-time processing of high-velocity.
- Absence of effective systems to manage heterogeneous urban datasets.
- Data privacy and security issues have not been solved.
- Scalability problems in extensive smart city implementations.

III. RESEARCH GAP

Even though big data analytics technologies are moving quickly in smart city applications, there are still some im-

TABLE I: Comparison of Existing Smart City Analytics Approaches

Paper	Domain	Technique	Advantages	Limitations
[1]	Urban	Data models	Improves planning	No real-time support
[2]	Big Data	Survey	Detailed study	No implementation
[3]	Transport	Prediction	Reduces congestion	Single-domain focus
[4]	Energy	Smart grid	Efficient usage	Scalability issues
[5]	Waste	IoT	Route optimization	Integration issues
[6]	Safety	Mining	Risk detection	Privacy concerns
[7]	Cloud	Distributed	Low latency	Complex design
[8]	AI/ML	ML models	High accuracy	Needs large data

portant research problems that need to be solved. Recent research underscores the necessity for more cohesive, scalable, secure, and sophisticated analytics frameworks tailored for contemporary smart city settings.

A. Lack of Integrated Multi-Domain Data Frameworks

Many existing studies focus on specific smart city domains such as transportation, energy, healthcare, or waste management independently. However, modern smart cities require a unified system capable of integrating heterogeneous data from multiple urban services simultaneously.

Prior studies have underscored the significance of integrated smart city architectures and the necessity for interoperability among urban systems [11], [21]. However, today's smart city implementations still operate in separate silos, which makes it harder to share data across domains and reduces the effectiveness of data-driven decision-making in cities [22]. Consequently, a robust multi-domain data integration framework is imperative.

B. Scalability Issues in Real-Time Data Processing

Smart city infrastructures generate massive volumes of real-time data through IoT devices, sensors, surveillance systems, and digital services. Processing these high-speed data streams efficiently remains a major technical challenge.

Existing studies highlight that traditional cloud-based systems struggle to manage the volume, velocity, and variety of smart city data [6], [21]. Recent research also emphasizes the need for edge computing and distributed architectures to enable real-time analytics in smart city environments [24]. Thus, scalable real-time data processing continues to be an unresolved research challenge.

C. Data Privacy and Security Concerns

When you collect and look at a lot of city data, it raises serious privacy and security concerns. Smart city platforms often deal with private information about people, infrastructure, and public services. Recent research underscores the increasing demand for privacy-preserving data sharing and secure analytics frameworks [25]. Nonetheless, numerous current solutions continue to emphasise performance and analytical functionalities over privacy safeguarding [22]. This results in a significant research deficiency in the advancement of secure

and privacy-conscious smart city analytics systems.

D. Limited Real-World Evaluation of Integrated Smart City Systems

Although numerous smart city frameworks have been proposed, many studies rely on theoretical models, simulations, or small-scale experiments. There is limited empirical validation of integrated analytics systems in real-world urban environments.

Research surveys emphasize the lack of large-scale performance evaluation of smart city analytics platforms [18], [21]. Comprehensive real-world experimentation is necessary to understand how integrated analytics systems can improve urban service efficiency and decision-making.

E. Need for Advanced Predictive and Decision-Support Models

Another major research gap is the limited development of advanced predictive and decision-support models capable of utilizing multi-domain urban datasets.

Artificial intelligence and machine learning have shown great potential in smart city applications [23]. However, existing systems still lack advanced predictive models that can transform raw urban data into actionable insights for proactive city management [22]. Developing intelligent decision-support systems remains a critical research direction.

Summary of Research Gap: The literature indicates that contemporary smart city research is deficient in a cohesive multi-domain analytics framework that is scalable, ensures privacy, and facilitates real-time predictive decision support. To make next-generation smart city platforms that are more sustainable, efficient, and improve quality of life, we need to deal with these problems.

IV. RESEARCH OBJECTIVES

The primary objectives of the research will be to:

- Study how various urban datasets can be effectively combined in a smart city scenario.
- Study real-time big data analytics approaches for monitoring urban systems.
- Study advanced predictive analytics approaches that can improve urban infrastructures and services.
- Study privacy-preserving approaches to make the storage of urban data secure.
- Study scalability and efficiency of the analytics platforms for large-scale smart cities' scenarios.
- Study IoT devices and sensor networks in building real-time urban datasets.
- Design a generic analytics platform capable of processing multidomain smart city datasets.
- Study how insights drawn from analysis of data can lead to smarter decision making in urban service management.

V. RESEARCH METHODOLOGY

The proposed methodology analyses and optimises urban infrastructure using a real-time multi-domain big data analytics

framework comprising data collection, preprocessing, integration, real-time processing, machine-learning-based modelling, prediction, and decision support.

A. System Architecture

The system is designed to operate in real time across distributed computing environments. Apache Kafka is used to handle large volumes of data streaming simultaneously from IoT devices. Apache Spark Structured Streaming then processes this data in batches, where it performs filtering, aggregation, and pattern detection. The processed data is stored in HDFS to support large-scale storage requirements, and it is also stored in MongoDB to enable efficient querying. Three Scikit-learn models are used to generate predictions related to traffic, energy, and environmental domains. The results are visualized through dashboards, and city administrators receive alerts through the system. Overall, Apache Kafka is used for data ingestion, Apache Spark for data processing, HDFS and MongoDB for storage, and Scikit-learn models for predictive analytics.




Fig. 1: Real-time smart city system architecture with ML model breakdown.

B. Implementation Details

The system is made using Python version 3 and Apache Spark. We use Kafka to make streams of data about traffic how much energy we use and what is happening in the environment. The Spark Structured Streaming piece looks at each group of data and does a few things to it. It filters out what we do not need, it adds up the data over a period of time and it looks for

things that are not normal. We have three Scikit-learn models that we train when the system is idle. Then we use them for predictions when the system is in use.

- **Random Forest** – using vehicle flow, road density, time of day and weather to predict traffic congestion
- **Linear Regression** – predicting energy consumption based on previous usage, temperature, occupancy.
- **Isolation Forest** – to detect anomalies in the environment using AQI, temperature, humidity and pollutant..

C. Algorithm for Real-Time Processing

In the real-time processing paradigm, it operates as a continuous flow of operations that ingests, analyses and processes urban data from diverse domains with minimal delay. The three kafka brokers have three topics for traffic, energy and environment and these topics are consumed with Spark Structured Streaming where the input is read as records in Distributed Dataframes. Each of these micro-batches will be preprocessed (nulls removal, min-max normalisation and time features extraction) and data will be aggregated along a window of five minutes. The models used for the 3 tasks are Random Forest, Linear Regression and Isolation Forest from Scikit-learn.

Algorithm 1 Real-Time Multi-Domain Urban Data Processing

- 1: Initialise Kafka brokers; create topics: *traffic*, *energy*, *environment*
- 2: Start Kafka producers to simulate IoT sensor streams
- 3: Configure Spark Structured Streaming session
- 4: Read stream from Kafka into Spark DataFrames
- 5: **Preprocessing:**
- 6: Remove null / out-of-range sensor values
- 7: Apply min-max normalisation to features
- 8: Extract time features (hour, day-of-week, season)
- 9: Perform windowed aggregation (5-min tumbling windows)
- 10: **Load pre-trained models:**
- 11: $M_{\text{traffic}} \leftarrow \text{Random Forest}$
- 12: $M_{\text{energy}} \leftarrow \text{Linear Regression}$
- 13: $M_{\text{env}} \leftarrow \text{Isolation Forest}$
- 14: **Inference on each micro-batch:**
- 15: $\hat{y}_{\text{traffic}} \leftarrow M_{\text{traffic}}.\text{predict}(X_{\text{traffic}})$
- 16: $\hat{y}_{\text{energy}} \leftarrow M_{\text{energy}}.\text{predict}(X_{\text{energy}})$
- 17: $\hat{y}_{\text{env}} \leftarrow M_{\text{env}}.\text{predict}(X_{\text{env}})$
- 18: Persist results to HDFS and MongoDB
- 19: **if** any $\hat{y}$ exceeds threshold τ **then**
- 20: Generate alert; push to dashboard
- 21: **end if**
- 22: Update dashboard KPIs; repeat from Step 4

D. Flowchart of the Proposed System

The detailed pipeline from raw sensor data to actionable insights is shown in Fig. 2. Each stage is marked with the technology or algorithm used.




Fig. 2: Detailed flowchart of the proposed real-time smart city analytics system.

E. Data Collection

The data is collected through IoT sensors, traffic surveillance cameras, environmental monitoring stations (AQI, temperature, humidity, pollution), intelligent energy metres, and public service platforms. The result is a significant amount of heterogeneous data which reflects the urban activity.

F. Data Preprocessing and Integration

The collected data is pre-processed to remove missing values, sensor failures and discrepancies. The numerical attributes are normalised using min-max scaling and time-related attributes (day-of-week, hour-of-day, seasonality) are derived to increase prediction accuracy. Apache Kafka separates the producer and consumer layers whereas Hadoop and Spark provide distributed computing and fault-tolerant storage.

G. Data Analysis and Modelling

Three models related to specific domains are trained off-line using Scikit-learn and saved for fast inference during on-line processing in the Spark workflow:

- 1) **Random Forest (traffic):** Accepts vehicle flow rate, road density, day-of-week, and weather attributes and predicts the congestion level.
- 2) **Linear Regression (energy):** Predicts hourly kWh demand based on historical consumption, ambient temperature, and

occupancy.

- 3) **Isolation Forest (environment):** Unsupervised anomaly detection using AQI, temperature, humidity, and pollution readings.

H. Optimisation and Decision Support

Model predictions feed optimisation strategies such as adaptive traffic-signal control and demand-response energy scheduling. A decision-support layer surfaces predictions, anomalies, and alerts through interactive dashboards, enabling city administrators to take timely, data-driven actions.

I. Performance Evaluation

The system is evaluated using the metrics summarised in Table II.

TABLE II: Evaluation Metrics for System Assessment

Metric	Domain	Description
Accuracy / RMSE	Traffic, Energy	Measures prediction accuracy of machine learning models
F1 Score	Environment	Evaluates precision and recall for anomaly detection
Processing Time (ms)	All Domains	Average latency per streaming micro-batch
Throughput (events/s)	All Domains	Number of sensor records processed per second
Scalability	System Level	Ability to handle increasing data volume efficiently
Response Time (ms)	Alert System	Time from threshold breach to alert notification

VI. EXPERIMENTAL SETUP

A. Tools and Technologies

The proposed architecture makes use of the following software tools and technologies that are best suited to handle real-time large-scale processing of urban big data:

- **Python 3.10** — Primary programming language for data pre-processing, feature extraction, machine learning model training and evaluation. The main Python packages used in the application are Pandas for data manipulation, NumPy for numerical computations, and Scikit-Learn for machine learning.
- **Apache Kafka 3.x** — Distributed messaging platform for handling real-time streaming data. Separate Kafka topics for traffic, energy, and environment data streams are maintained.
- **Apache Spark 3.x (Structured Streaming)** — Distributed computing engine for real-time streaming analytics such as micro-batch processing, window aggregation and filtering, and real-time inference. Spark DataFrames are used for parallel processing of each micro-batch in a distributed cluster environment.
- **Hadoop HDFS** — Distributed file storage system for fault-tolerant long-term archiving of processed records and model

outputs.

- **MongoDB** — NoSQL document database for real-time access and retrieval of real-time inference results by the dashboard component.
- **Scikit-learn 1.3** — Machine learning package used to train, evaluate and serialize the Random Forest, Linear Regression, and Isolation Forest models for real-time inference in the Spark stream processing pipeline.
- **Jupyter Notebook** — Main platform for exploratory analysis, experimentation, and visualizations during the development phase.

B. Dataset Description

Three publicly available datasets are utilized for the experiments in this study, one per urban domain. They provide around 86,000 instances regarding traffic, energy, and environment. All datasets are obtained from Kaggle and the UCI Machine Learning Repository.

- **Traffic Domain — Metro Interstate Traffic Volume Dataset [34]:** Contains 48,204 hourly traffic volume records from Interstate 94, Minnesota, USA, including weather condition, temperature, rain, snow, and time-of-day features.
- **Energy Domain — Smart Home Energy Consumption Dataset [35]:** Contains hourly household appliance energy consumption readings with temperature, humidity, and time-of-day features, collected from a real smart home deployment.
- **Environmental Domain — Air Quality UCI Dataset [36]:** Contains 9,358 hourly sensor readings from an Italian city including CO, NOx, NO2, benzene concentration, temperature, and relative humidity, recorded from March 2004 to February 2005.

All the three data sets are separately preprocessed to remove any missing values, normalize numeric columns via min-max normalization, and extract time-oriented attributes, hour of the day, day of the week, and season.

C. Evaluation Methodology

Each model will be evaluated using appropriate performance metrics based on the type of task it performs. For the supervised models, the performance will be determined on an 80/20 train/test splitting of the datasets. However, the unsupervised model will be tested on anomalies in a separate dataset that has been labelled manually.

- **Random Forest (Traffic — Classification):** Evaluated by Accuracy, Precision, Recall and F1 Score based on the holdout 20% test split of the traffic dataset.
- **Linear Regression (Energy — Regression):** Evaluated using RMSE, MAE, and the coefficient of determination (R^2) based on the holdout 20% test split of the energy dataset.
- **Isolation Forest (Environment — Anomaly Detection):** Evaluated using Precision, Recall, and F1 Score against a set of anomalies from a manual labelling of the environmental dataset. The contamination factor will be set at 0.05.

System-wise, the entire system pipeline is evaluated using three more dimensions. Latency will refer to the total time

taken for data to move from ingestion into Kafka to the generation of alerts at the dashboard interface. Throughput is defined as the rate at which records are processed in a second by the Spark Streaming engine. Scalability refers to how the processing time varies as the amount of input data is increased gradually.

D. Parameter Settings

The system and model parameters are carefully configured to ensure efficient real-time processing and accurate prediction across all three domains. Table III summarises the key experimental parameters used in this study.

- **Random Forest** is configured with 100 decision trees and a maximum depth of 10. The model is trained on 80% of the traffic dataset (38,563 records) and tested on the remaining 20% (9,641 records).
- **Linear Regression** uses default Scikit-learn settings with no regularisation, trained and tested on an 80/20 split of the energy dataset (23,258 training, 5,815 test records).
- **Isolation Forest** is configured with a contamination factor of 0.05, meaning approximately 468 of the 9,358 environmental records are expected to be anomalous.
- **Spark Structured Streaming** is configured with a micro-batch trigger interval of 5 seconds and a tumbling aggregation window of 5 minutes.
- **Apache Kafka** is configured with a replication factor of 1 and a single partition per topic, appropriate for the simulated single-node experimental environment.

TABLE III: Experimental Parameter Settings

Parameter	Value
Programming Language	Python 3.10
ML Library	Scikit-learn 1.3
Traffic Dataset	Metro Interstate (48,204 records)
Energy Dataset	Smart Home Consumption (29,073 records)
Environment Dataset	Air Quality UCI (9,358 records)
Train / Test Split	80% / 20%
Streaming Tool	Apache Kafka 3.x
Processing Tool	Apache Spark 3.x
Micro-batch Interval	5 seconds
Aggregation Window	5-minute tumbling window
Storage	HDFS + MongoDB
RF Trees / Max Depth	100 / 10
Isolation Forest Contam.	0.05 (5%)
Traffic Metrics	Accuracy, Precision, Recall, F1
Energy Metrics	RMSE, MAE, R^2
Env. Metrics	Precision, Recall, F1
System Metrics	Latency, Throughput, Scalability

VII. RESULTS AND DISCUSSION

The proposed real-time multi-domain smart city analytics framework is evaluated across three urban domains — traffic, energy, and environment — using the datasets and parameters described in Section IV. Results are reported for individual

ML model performance, system-level pipeline metrics, and a comparative analysis against existing approaches.

A. ML Model Performance

Each machine learning model is evaluated independently on its held-out 20% test set using domain-appropriate metrics.

1) Traffic Prediction — Random Forest

The Random Forest classifier achieves strong performance on the traffic congestion classification task. Table IV summarises the classification results across the three congestion levels.

TABLE IV: Random Forest — Traffic Classification Results

Class	Precision	Recall	F1	Support
Low Congestion	0.94	0.96	0.95	3,842
Medium Congestion	0.91	0.89	0.90	3,201
High Congestion	0.93	0.92	0.92	2,598
Overall	0.93	0.92	0.92	9,641

The overall accuracy of the Random Forest model is **92.4%**, with a weighted F1 Score of 0.92, demonstrating reliable congestion level prediction across all three classes.

2) Energy Forecasting — Linear Regression

Energy Demand Prediction — Linear Regression Model Performance The Linear Regression model's performance in predicting energy demand per hour is provided in Table V.

TABLE V: Linear Regression — Energy Forecasting Results

Metric	Value
Root Mean Square Error (RMSE)	0.847 kWh
Mean Absolute Error (MAE)	0.623 kWh
Coefficient of Determination (R^2)	0.883
Mean Absolute Percentage Error (MAPE)	8.2%

With the value of R^2 equal to 0.883, the model captures 88.3% of the variation in energy use, demonstrating high accuracy and efficiency in energy forecasting, which can be used in real-time energy demand management.

3) Environmental Anomaly Detection — Isolation Forest

The Isolation Forest model is assessed using labelled anomaly data with contamination of 0.05. It identifies 468 of 9,358 total environmental records as anomalies.

TABLE VI: Isolation Forest — Anomaly Detection Results

Metric	Value
Precision	0.89
Recall	0.86
F1 Score	0.87
False Positive Rate	4.8%
Anomalies Detected	447 / 468

Out of the detected anomalies, a total of 447 were verified as correct, demonstrating that the system achieved a very low error rate of only 4.8%.

B. Graph 1: Model Accuracy Comparison

Fig. 3 shows a comparison between the accuracy and F1 scores for the three proposed machine learning models and their corresponding baseline counterparts for the domains of traffic, energy, and environment, which are Naive Bayes, simple moving average, and statistical thresholding, respectively. Different colors have been used to distinctly differentiate the domains from one another.




Fig. 3: Accuracy and F1 score comparison between baseline methods and proposed ML models across all three urban domains. Orange (hatched) = baseline; Blue (solid) = proposed.

The proposed models beat the traditional approaches on average by **15.3%**. This improvement is maximum in case of traffic prediction (+14%). It occurs because Random Forest algorithm can predict the non-linear dependency between the variables of weather condition, time of the day, and traffic density while Naive Bayes algorithm cannot do it due to conditional independence hypothesis. Energy forecast has been improved on +14% since Linear Regression algorithm can detect highly significant linear dependency between temperature, occupancy, and demand for kWh per hour that can not be predicted by Simple Moving Average. Environmental anomaly detection has seen the greatest improvement +16%. It happens because Isolation Forest can detect multivariate anomalies among air quality index, humidity, and pollution while threshold based algorithms miss those.

C. Graph 2: Processing Latency Comparison

Fig. 4 shows the end-to-end processing latency — measured from Kafka ingestion to dashboard alert generation — for the proposed streaming system versus a traditional batch processing baseline, across five increasing dataset sizes.




Fig. 4: End-to-end latency comparison between batch processing (red dashed) and the proposed real-time streaming pipeline (teal solid) as dataset size increases.

The proposed system maintains an average latency of **99.2 ms** compared to **491.8 ms** for batch processing, achieving a **79.8%** latency reduction. The near-flat slope of the proposed curve confirms that Spark Structured Streaming absorbs increasing data volumes with minimal latency growth, while batch processing latency rises steeply.

D. Graph 3: System Throughput vs Worker Nodes

Fig. 5 illustrates the throughput of the proposed system — measured in records processed per second — as the number of distributed Apache Spark worker nodes scales from 1 to 4. This directly demonstrates the horizontal scalability of the framework.




Fig. 5: System throughput (records/second) vs. number of Spark worker nodes. Violet bars = actual throughput achieved; green dashed = ideal linear scaling reference.

Throughput ranges from 3,200 records/second for a single node to 11,400 records/second for four nodes. The observed throughput is closely aligned with the ideal scaling of linear form, proving near-linear horizontal scalability. The minimal

difference between ideal and observed throughput results from inter-node communication overhead in Spark clusters.

E. Graph 4: Scalability — Performance vs Data Volume

Figure 6 shows a comparison between the performance score of the proposed model and the existing single domain batch processing system based on benchmarking the performance score with increasing the number of records from 10,000 to 50,000. The performance score of both systems is evaluated on a weighted basis for prediction accuracy, latency, and throughput and is plotted between 0 to 100 percent. It can be seen that the performance score for the bench marked existing single domain batch processing system decreases consistently because of its centralized and non-distributed approach. On the other hand, the performance score of the proposed model increases from 91% to 95% by utilizing Apache Spark’s distributed architecture and parallel computing capability that allows efficient distribution of data among worker nodes without overloading any particular node.




Fig. 6: Scalability comparison: system performance score vs. increasing data volume. Orange dashed = existing system (degrading); Blue solid = proposed system (stable).

The current system deteriorates from 82% to 55% efficiency as data volume increases, indicating the weaknesses of the centralized batch processing model. The suggested system remains consistent at 91% efficiency and shows a slight improvement to 95%, proving that the distributed Spark processing model successfully sustains efficiency as the data volume of smart cities increases.

F. Graph 5: Domain-wise F1 Score Breakdown

Fig. 7 shows a more detailed analysis of the precision, recall, and F1 score for each of the three ML algorithms, which allows one to gain an understanding of the performance of individual models.




Fig. 7: Precision, Recall, and F1 Score breakdown for each ML model across the three urban domains. Blue = Precision, Green = Recall, Red = F1 Score.

All three ML models demonstrate well-balanced precision and recall, which means that there is no dominance of either type of errors. The traffic dataset Random Forest model obtains the best results because it works with the largest and richest dataset of all three.

G. Comparison with Existing Systems

Table VII provides a side-by-side comparison between our approach and others used for smart city analytics.

TABLE VII: Comparison with Existing Systems

Feature	Existing Systems	Proposed System
Data Handling	Single-domain	Multi-domain (3 domains)
Processing	Batch processing	Real-time streaming
Latency	~490 ms	~99 ms
Accuracy	71–78%	87–92%
Scalability	Degrades with volume	Near-linear scaling
ML Models	Single general model	Domain-specific models
Anomaly Detection	Rule-based thresholds	Isolation Forest (ML)
Privacy Support	None	Anonymisation + TLS
Decision Support	Basic dashboards	Real-time alerts + KPIs

H. Discussion

These results confirm that the proposed approach addresses the main shortcomings of the existing smart city analytics systems in all three analyzed dimensions.

Employing specialized ML models for a particular industry leads to significantly more precise forecasts than relying on just one general purpose ML model. Random Forest detects nonlinear relationships between traffic indicators, Linear Regression models the nearly linear behavior of energy consumption, and Isolation Forest detects anomalies in environmental conditions without requiring labeled training data.

The Spark Structured Streaming pipeline cuts processing latency by 79.8% compared to batch processing, and adding worker nodes makes throughput grow almost linearly. This shows that real-time distributed processing is both possible and very efficient for smart city deployments that need quick responses.

The multi-domain integration approach enables cross-domain insights unavailable to single-domain systems. Correlating traffic congestion spikes with energy surges and air quality degradation provides city administrators with a comprehensive, unified view of urban conditions for more proactive decision-making.

A limitation of the current study is that results are based on simulated streaming using historical datasets rather than live IoT deployments. Future work will focus on live testbed deployment, federated learning for privacy-preserving distributed model training, and extension to additional urban domains such as waste management and public safety.

VIII. CONCLUSION

The modern cities require complex systems to provide critical services like transportation, energy, and environment. Traditional single domain batch-processing systems lack the ability to handle the real-time processing capabilities, which results in poor decision-making and scaling of resources. In this paper, we presented a new multi-domain big data real-time processing approach utilizing distributed stream processing, machine learning, and privacy-preserving techniques.

Our solution combines three Apache open-source tools including Apache Kafka, Apache Spark Structured Streaming, and three Scikit-learn algorithms namely Random Forest for traffic prediction, Linear Regression for energy forecasting, and Isolation Forest for environmental anomaly detection.

The experimentally achieved results include 92.4% overall accuracy in predicting traffic, an R^2 coefficient of 0.883 for energy forecasting, and an environmental anomaly detection F1-score of 0.87. From the perspective of overall system performance, our system achieves a latency reduction of 79.8% compared with traditional batch processing, with almost linearly scaled throughput. We ensure data privacy through data anonymization, TLS encryption, differential privacy, and Role-Based Access Control (RBAC) mechanisms.

In summary, the presented architecture illustrates the potential for integrating real-time multi-domain data, domain-specific machine learning, and privacy protection into a highly effective smart city analysis system. The future scope of work is expected to include the deployment of IoT testbeds, federated learning techniques, and applications within other domains like waste disposal and public security.

IX. FUTURE WORK

While the proposed framework demonstrates strong performance across traffic, energy, and environmental domains, several directions remain open for further research and development.

Live IoT Deployment: The current system is evaluated using historical datasets simulated as real-time streams. Future

work will focus on deploying the framework on a live smart city testbed with real IoT sensor infrastructure, enabling evaluation under actual urban conditions including sensor noise, network failures, and unpredictable data spikes.

Federated Learning: To further strengthen privacy, federated learning will be explored as an alternative to centralised model training. This approach allows ML models to be trained across multiple distributed city zones without raw data ever leaving the source node, significantly reducing privacy risk while maintaining prediction accuracy.

Edge Computing Integration: Integrating edge computing nodes into the pipeline will enable data preprocessing and inference to occur closer to the data source, reducing network congestion and further lowering end-to-end latency beyond what the current cloud-based Spark deployment achieves. This is particularly relevant for time-critical applications such as emergency response and accident detection.

5G Network Support: The adoption of 5G networks will significantly enhance the data ingestion capacity of the Kafka layer, enabling support for a much larger number of concurrent IoT devices and higher-frequency sensor sampling rates across all urban domains.

Advanced Privacy Mechanisms: Future work will look into secure multi-party computation and blockchain for data provenance, which can serve as privacy enhancing techniques, allowing independent city departments to share their data securely across domains with complete traceability of data access.

Domain Expansion: The proposed framework will be applied to other domains of urban administration such as garbage disposal, water distribution networks, and public safety response optimization, thus further realizing the concept of an integrated smart city analytics framework.

Intelligent Decision Support: Intelligent visualization features and recommender systems will be embedded in future dashboard frameworks to help city administrators make informed decisions during urban planning and policy making.

ACKNOWLEDGMENT

The author would like to express sincere gratitude to the faculty members of the Department of Computer Science, Chitkara University Institute of Engineering and Technology, for their valuable guidance and academic support during the preparation of this research work.

The author also acknowledges the contributions of researchers and scholars in the fields of big data analytics, Internet of Things, and smart city technologies, whose published work provided important insights and references for this study.

REFERENCES

- [1] M. Batty, "Big data, smart cities and city planning," *Dialogues in Human Geography*, vol. 3, no. 3, pp. 274–279, 2013.
- [2] M. Chen, S. Mao, and Y. Liu, "Big data: A survey," *Mobile Networks and Applications*, vol. 19, no. 2, pp. 171–209, 2014.
- [3] J. Gubbi, R. Buyya, S. Marusic, and M. Palaniswami, "Internet of Things (IoT): A vision, architectural elements, and future directions," *Future Generation Computer Systems*, vol. 29, no. 7, pp. 1645–1660, 2013.
- [4] S. Zanella *et al.*, "Internet of Things for smart cities," *IEEE Internet of Things Journal*, vol. 1, no. 1, pp. 22–32, 2014.

- [5] R. Kitchin, "The real-time city? Big data and smart urbanism," *Geo-Journal*, vol. 79, no. 1, pp. 1–14, 2014.
- [6] I. Hashem *et al.*, "The role of big data in smart city," *International Journal of Information Management*, vol. 36, no. 5, pp. 748–758, 2016.
- [7] A. Botta *et al.*, "Integration of cloud computing and Internet of Things: A survey," *Future Generation Computer Systems*, vol. 56, pp. 684–700, 2016.
- [8] Y. Zheng, "Urban computing: Concepts, methodologies, and applications," *ACM Transactions on Intelligent Systems and Technology*, vol. 5, no. 3, pp. 1–55, 2014.
- [9] J. Jin *et al.*, "An information framework for creating a smart city through IoT," *IEEE Internet of Things Journal*, vol. 1, no. 2, pp. 112–121, 2014.
- [10] L. Anthopoulos, *Smart City Implementation*, Springer, 2017.
- [11] H. Chourabi *et al.*, "Understanding smart cities: An integrative framework," in *Proc. HICSS*, 2012, pp. 2289–2297.
- [12] R. G. Hollands, "Will the real smart city please stand up?" *City*, vol. 12, no. 3, pp. 303–320, 2008.
- [13] T. Nam and T. A. Pardo, "Conceptualizing smart city," in *Proc. Digital Government Research*, 2011, pp. 282–291.
- [14] C. Harrison and I. A. Donnelly, "A theory of smart cities," in *Proc. ISSS*, 2011.
- [15] V. Albino *et al.*, "Smart cities: Definitions and performance," *Journal of Urban Technology*, vol. 22, no. 1, pp. 3–21, 2015.
- [16] A. Caragliu *et al.*, "Smart cities in Europe," *Journal of Urban Technology*, vol. 18, no. 2, pp. 65–82, 2011.
- [17] K. K. Patel and S. M. Patel, "Internet of Things-IOT," *IJESCI*, vol. 6, no. 5, pp. 6122–6131, 2016.
- [18] M. Al Nuaimi *et al.*, "Applications of big data to smart cities," *JISA*, vol. 6, no. 1, pp. 1–15, 2015.
- [19] S. Yin and O. Kaynak, "Big data for modern industry," *Proceedings of the IEEE*, vol. 103, no. 2, pp. 143–146, 2015.
- [20] J. Manyika *et al.*, "Big data: The next frontier," McKinsey, 2011.
- [21] X. Zhang, Z. Chen, and Y. Li, "Smart city data analytics: A review," *IEEE Access*, vol. 8, pp. 11321–11345, 2020.
- [22] A. K. Singh and R. R. Sharma, "Big data analytics for smart cities: Applications and challenges," *Sustainable Cities and Society*, vol. 61, 2020.
- [23] M. A. Jan *et al.*, "Artificial intelligence for smart cities," *IEEE Communications Magazine*, vol. 58, no. 1, pp. 96–102, 2021.
- [24] S. Sharma, S. Chen, and J. Park, "Edge computing in smart cities," *IEEE Network*, vol. 35, no. 3, pp. 92–98, 2021.
- [25] Y. Liu *et al.*, "Privacy-preserving data sharing in smart cities," *IEEE Internet of Things Journal*, vol. 9, no. 4, pp. 2345–2356, 2022.
- [26] H. Wang, J. Li, and Q. Zhao, "Real-time urban traffic prediction using deep learning and Apache Spark streaming," *IEEE Transactions on Intelligent Transportation Systems*, vol. 24, no. 3, pp. 3421–3435, 2023.
- [27] F. Al-Turjman, M. Nawaz, and U. Rohit, "Intelligence in IoT-enabled smart cities: Federated learning, green computing, and deep learning," *IEEE Access*, vol. 11, pp. 14933–14950, 2023.
- [28] R. Minerva, G. M. Lee, and N. Crespi, "Digital twin in the IoT context: A survey on technical features, scenarios, and architectural models," *Proceedings of the IEEE*, vol. 111, no. 3, pp. 201–227, 2023.
- [29] Z. Chen, X. Wang, and Y. Zhang, "Privacy-preserving federated learning for smart city applications," *IEEE Internet of Things Journal*, vol. 10, no. 8, pp. 6891–6903, 2023.
- [30] A. Masood *et al.*, "Federated learning-based anomaly detection for IoT security in smart cities," *IEEE Communications Magazine*, vol. 61, no. 2, pp. 72–78, 2023.
- [31] S. Koppu, V. M. Viswanatham, and P. K. Gupta, "Energy consumption forecasting in smart grids using machine learning: A survey," *Sustainable Cities and Society*, vol. 98, pp. 104–118, 2024.
- [32] M. Raza *et al.*, "Real-time air quality monitoring and anomaly detection using machine learning in smart city environments," *Environmental Science and Pollution Research*, vol. 31, no. 1, pp. 1045–1062, 2024.
- [33] T. Zhao, L. Chen, and W. Liu, "Scalable big data analytics for smart city infrastructure using Apache Spark and Kafka: Architecture and performance evaluation," *Future Generation Computer Systems*, vol. 152, pp. 312–327, 2024.
- [34] A. Tanwar, "Metro Interstate Traffic Volume Dataset," Kaggle, 2019. [Online]. Available: <https://www.kaggle.com/datasets/anshtanwar/metro-interstate-traffic-volume>. Accessed: Apr. 2026.
- [35] Maxwell, "Smart Home Energy Consumption Dataset," Kaggle, 2024. [Online]. Available: <https://www.kaggle.com/datasets/mexwell/smart-home-energy-consumption>. Accessed: Apr. 2026.
- [36] S. De Vito, E. Massera, G. Piga, L. Martinotto, and G. Di Francia, "Air Quality Dataset," UCI Machine Learning Repository, 2008. [Online]. Available: <https://www.kaggle.com/datasets/fedesoriano/air-quality-data-set>. Accessed: Apr. 2026.